

# Credit Card Approval Prediction Using Sklearn

## Table of Content

- [1 Feature Engineering](#)
  - [1.1 Response Variable](#)
  - [1.2 Features](#)
    - [1.2.1 Binary Features](#)
      - [1.2.1.1 Gender](#)
    - [1.2.3 Categorical Features](#)
  - [2 Algorithms](#)
    - [2.1 Logistic Regression](#)
    - [2.2 Decision Tree](#)
    - [2.3 Random Forest](#)
    - [2.4 SVM](#)
    - [2.5 LightGBM](#)
    - [2.6 Xgboost](#)
    - [2.7 Keras Neural Networks](#)

```
In [1]: # %matplotlib inline
# %config InlineBackend.figure_format = 'svg'

import warnings
warnings.filterwarnings('ignore')
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
from imblearn.over_sampling import SMOTE
import itertools

from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, confusion_matrix
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier

from xgboost import XGBClassifier
from lightgbm import LGBMClassifier
from catboost import CatBoostClassifier
from sklearn import svm
from sklearn.ensemble import RandomForestClassifier
```

Using TensorFlow backend.

```
In [2]: # plt.rcParams['figure.facecolor'] = 'white'
```

## Binary Features

```

In [3]: # Calculate information value
def calc_iv(df, feature, target, pr=False):
    lst = []
    df[feature] = df[feature].fillna("NULL")

    for i in range(df[feature].nunique()):
        val = list(df[feature].unique())[i]
        lst.append([feature,                                     # Variable
                    val,                                         # Value
                    df[df[feature] == val].count()[feature],    # All
                    df[(df[feature] == val) & (df[target] == 0)].count()[feature], # Good (think: Fraud == 0)
                    df[(df[feature] == val) & (df[target] == 1)].count()[feature]]) # Bad (think: Fraud == 1)

    data = pd.DataFrame(lst, columns=['Variable', 'Value', 'All', 'Good', 'Bad'])
    data['Share'] = data['All'] / data['All'].sum()
    data['Bad Rate'] = data['Bad'] / data['All']
    data['Distribution Good'] = (data['All'] - data['Bad']) / (data['All'].sum() - data['Bad'].sum())
    data['Distribution Bad'] = data['Bad'] / data['Bad'].sum()
    data['WoE'] = np.log(data['Distribution Good'] / data['Distribution Bad'])

    data = data.replace({'WoE': {np.inf: 0, -np.inf: 0}})

    data['IV'] = data['WoE'] * (data['Distribution Good'] - data['Distribution Bad'])

    data = data.sort_values(by=['Variable', 'Value'], ascending=[True, True])
    data.index = range(len(data.index))

    if pr:
        print(data)
        print('IV = ', data['IV'].sum())

    iv = data['IV'].sum()
    print('This variable\'s IV is:', iv)
    print(df[feature].value_counts())
    return iv, data

def convert_dummy(df, feature, rank=0):
    pos = pd.get_dummies(df[feature], prefix=feature)
    mode = df[feature].value_counts().index[rank]
    biggest = feature + '_' + str(mode)
    pos.drop([biggest], axis=1, inplace=True)

```

```
df.drop([feature],axis=1,inplace=True)
df=df.join(pos)
return df

def get_category(df, col, binsnum, labels, qcut = False):
    if qcut:
        localdf = pd.qcut(df[col], q = binsnum, labels = labels) # quantile cut
    else:
        localdf = pd.cut(df[col], bins = binsnum, labels = labels) # equal-length cut

    localdf = pd.DataFrame(localdf)
    name = 'gp' + '_' + col
    localdf[name] = localdf[col]
    df = df.join(localdf[name])
    df[name] = df[name].astype(object)
    return df

def plot_confusion_matrix(cm, classes,
                           normalize=False,
                           title='Confusion matrix',
                           cmap=plt.cm.Blues):

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    print(cm)

    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()
    tick_marks = np.arange(len(classes))
    plt.xticks(tick_marks, classes)
    plt.yticks(tick_marks, classes)

    fmt = '.2f' if normalize else 'd'
    thresh = cm.max() / 2.
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        plt.text(j, i, format(cm[i, j], fmt),
                 horizontalalignment="center",
                 color="white" if cm[i, j] > thresh else "black")

    plt.tight_layout()
```

```
plt.ylabel('True label')  
plt.xlabel('Predicted label')
```

## Feature Engineering

```

In [4]: def return_no(x):
        if x==None:
            return "No"

def data_preprocessing(data):

    ## Feature Engineering
    # find all users' account open month.
    begin_month=pd.DataFrame(data.groupby(["ID"])["Vintage"].agg(min))
    begin_month=begin_month.rename(columns={'Vintage':'begin_month'})
    new_data=pd.merge(data,begin_month,how="left",on="ID") #merge to record data

    del begin_month

    ##Generally, users in risk should be in 3%, thus I choose users who overdue for more than 60 days as target risk u
sers. Those samples are marked as '1', else are '0'.
    new_data['target']=data['Is_Active']
    new_data.loc[new_data['target']=='Yes','target']=1
    new_data.loc[new_data['target']=='No','target']=0
    print(new_data['target'].value_counts())
    new_data['target'].value_counts(normalize=True)

    #features
    # new_data.dropna()
    # new_data = new_data.mask(new_data == 'NULL').dropna()

    ivtable=pd.DataFrame(new_data.columns,columns=['variable'])
    ivtable['IV']=None

    # for col in columns:
    #     try:
    # #         new_data[col].replace(np.NaN, new_data[col].mean())
    #         new_data[col].replace(None, new_data[col].mean())
    #     except:pass

    #No      144357, Yes      72043
    new_data['Credit_Product'] = new_data['Credit_Product'].apply(lambda x: return_no(x))

    from sklearn.preprocessing import LabelEncoder

```

```
label_encoder = LabelEncoder()

l_encoder = label_encoder.fit(new_data['Gender'])
new_data['Gender'] = l_encoder.transform(new_data['Gender'])

l_encoder = label_encoder.fit(new_data['Region_Code'])
new_data['Region_Code'] = l_encoder.transform(new_data['Region_Code'])

l_encoder = label_encoder.fit(new_data['Occupation'])
new_data['Occupation'] = l_encoder.transform(new_data['Occupation'])

l_encoder = label_encoder.fit(new_data['Channel_Code'])
new_data['Channel_Code'] = l_encoder.transform(new_data['Channel_Code'])

l_encoder = label_encoder.fit(new_data['Credit_Product'])
new_data['Credit_Product'] = l_encoder.transform(new_data['Credit_Product'])

#     l_encoder = label_encoder.fit(new_data['Is_Active'])
#     new_data['Is_Active'] = l_encoder.transform(new_data['Is_Active'])

del label_encoder, l_encoder

#     print(data['Is_Active'].value_counts())
#     data['Is_Active'].value_counts(normalize=True)

#Gender
print(new_data['Gender'].value_counts())
iv, data = calc_iv(new_data, 'Gender', 'target')
ivtable.loc[ivtable['variable']=='Gender', 'IV']=iv

#Avg_Account_Balance
#     new_data['Avg_Account_Balance']=new_data['Avg_Account_Balance'].astype(object)
#     new_data['Avg_Account_Balance'] = new_data['Avg_Account_Balance']/10000
new_data['Avg_Account_Balance'] = new_data['Avg_Account_Balance'].apply(lambda x: x/10000)

print(new_data['Avg_Account_Balance'].value_counts(bins=10,sort=False))
new_data['Avg_Account_Balance'].plot(kind='hist',bins=50,density=True)

return new_data
```





```
In [5]: new_data = data_preprocessing(pd.read_csv("../input/jobathon-may-2021/train.csv", encoding = 'utf-8') )  
  
print(new_data.columns)  
new_data.head()
```


```

0    150290
1     95435
Name: target, dtype: int64
1    134197
0    111528
Name: Gender, dtype: int64
This variable's IV is: 0.01743281354405051
1    134197
0    111528
Name: Gender, dtype: int64
(1.0450000000000002, 105.391]    148577
(105.391, 208.703]              73815
(208.703, 312.016]             15340
(312.016, 415.328]              4834
(415.328, 518.64]               1585
(518.64, 621.952]               738
(621.952, 725.264]              418
(725.264, 828.577]              282
(828.577, 931.889]              125
(931.889, 1035.201]             11
Name: Avg_Account_Balance, dtype: int64
Index(['ID', 'Gender', 'Age', 'Region_Code', 'Occupation', 'Channel_Code',
      'Vintage', 'Credit_Product', 'Avg_Account_Balance', 'Is_Active',
      'Is_Lead', 'begin_month', 'target'],
      dtype='object')

```

Out[5]:

	ID	Gender	Age	Region_Code	Occupation	Channel_Code	Vintage	Credit_Product	Avg_Account_Balance	Is_Active	Is_Lead	begin
0	NNVBBKZB	0	73	18	1	2	43	0	104.5696	No	0	
1	IDD62UNG	0	30	27	2	0	32	0	58.1988	No	0	
2	HD3DSEMC	0	56	18	3	2	26	0	148.4315	Yes	0	
3	BF3NC7KV	1	34	20	2	0	19	0	47.0454	No	0	
4	TEASRWXV	0	30	32	2	0	33	0	88.6787	No	0	



```
In [6]: test_data = data_preprocessing(pd.read_csv("../input/jobathon-may-2021/test.csv", encoding = 'utf-8') )  
test_data.head()
```


```

0    63797
1    41515
Name: target, dtype: int64
1    57705
0    47607
Name: Gender, dtype: int64
This variable's IV is: 0.020052086601170444
1    57705
0    47607
Name: Gender, dtype: int64
(1.27, 101.122]    60798
(101.122, 199.985]  33196
(199.985, 298.848]  7263
(298.848, 397.71]   2343
(397.71, 496.573]   872
(496.573, 595.435]   396
(595.435, 694.298]   201
(694.298, 793.161]   152
(793.161, 892.023]    63
(892.023, 990.886]    28
Name: Avg_Account_Balance, dtype: int64

```

Out[6]:

	ID	Gender	Age	Region_Code	Occupation	Channel_Code	Vintage	Credit_Product	Avg_Account_Balance	Is_Active	begin_month
0	VBENBARO	1	29	4	1	0	25	0	74.2366	No	25
1	CCMEWNKY	1	43	18	1	1	49	0	92.5537	No	49
2	VK3KGA9M	1	31	20	2	0	14	0	21.5949	No	14
3	TT8RPZVC	1	29	22	1	0	33	0	86.8070	No	33
4	SHQZEY TZ	0	29	20	1	0	19	0	65.7087	No	19



## Algorithms

```
In [7]: Y = new_data['Is_Lead']  
X = new_data[['Gender', 'Age', 'Region_Code', 'Occupation', 'Channel_Code',  
             'Credit_Product', 'Avg_Account_Balance', 'target', 'begin_month']]  
X.head()
```

Out[7]:

	Gender	Age	Region_Code	Occupation	Channel_Code	Credit_Product	Avg_Account_Balance	target	begin_month
0	0	73	18	1	2	0	104.5696	0	43
1	0	30	27	2	0	0	58.1988	0	32
2	0	56	18	3	2	0	148.4315	1	26
3	1	34	20	2	0	0	47.0454	0	19
4	0	30	32	2	0	0	88.6787	0	33

- Using Synthetic Minority Over-Sampling Technique( SMOTE ) to overcome sample imbalance problem.

```
In [8]: Y = Y.astype('int')
X_balance, Y_balance = SMOTE().fit_sample(X, Y)
X_balance = pd.DataFrame(X_balance, columns = X.columns)
```

- After over sampling, the number between 1 and 0 is balanced. It can be seen from the confusion matrix.

```
In [9]: X_train, X_test, y_train, y_test = train_test_split(X_balance, Y_balance, test_size=0.2,
                                                         random_state = 10086)
print(X_train.size, X_test.size, y_train.size, y_test.size)

2699091 674775 299899 74975
```

## Logistic Regression

$$\log\left(\frac{p}{1-p}\right) = \beta_0 + \beta_1 x_1 + \cdots + \beta_q x_q$$

```
In [10]: # model = LogisticRegression(C=0.8,
#                                     random_state=0,
#                                     solver='lbfgs')
# model.fit(X_train, y_train)
# y_predict = model.predict(X_test)

# print('Accuracy Score is {:.5}'.format(accuracy_score(y_test, y_predict)))
# print(pd.DataFrame(confusion_matrix(y_test, y_predict)))

# sns.set_style('white')
# class_names = ['0', '1']
# plot_confusion_matrix(confusion_matrix(y_test, y_predict),
#                       classes=class_names, normalize=True,
#                       title='Normalized Confusion Matrix: Logistic Regression')
```

## Decision Tree

```
In [11]: # model = DecisionTreeClassifier(max_depth=12,
#                                     min_samples_split=8,
#                                     random_state=1024)
# model.fit(X_train, y_train)
# y_predict = model.predict(X_test)

# print('Accuracy Score is {:.5}'.format(accuracy_score(y_test, y_predict)))
# print(pd.DataFrame(confusion_matrix(y_test, y_predict)))

# plot_confusion_matrix(confusion_matrix(y_test, y_predict),
#                       classes=class_names, normalize = True,
#                       title='Normalized Confusion Matrix: CART')
```

## Random Forest

```
In [12]: # model = RandomForestClassifier(n_estimators=250,
#                                     max_depth=12,
#                                     min_samples_leaf=16
#                                     )
# model.fit(X_train, y_train)
# y_predict = model.predict(X_test)

# print('Accuracy Score is {:.5}'.format(accuracy_score(y_test, y_predict)))
# print(pd.DataFrame(confusion_matrix(y_test, y_predict)))

# plot_confusion_matrix(confusion_matrix(y_test, y_predict),
#                       classes=class_names, normalize = True,
#                       title='Normalized Confusion Matrix: Ramdom Forests')
```

## SVM



```
In [13]: # model = svm.SVC(C = 0.8,
#               kernel='linear')
# model.fit(X_train, y_train)
# y_predict = model.predict(X_test)

# print('Accuracy Score is {:.5}'.format(accuracy_score(y_test, y_predict)))
# print(pd.DataFrame(confusion_matrix(y_test,y_predict)))

# plot_confusion_matrix(confusion_matrix(y_test,y_predict),
#               classes=class_names, normalize = True,
#               title='Normalized Confusion Matrix: SVM')
```

## LightGBM

```
In [14]: # model = LGBMClassifier(num_leaves=31,
#               max_depth=8,
#               learning_rate=0.02,
#               n_estimators=250,
#               subsample = 0.8,
#               colsample_bytree =0.8
#               )
# model.fit(X_train, y_train)
# y_predict = model.predict(X_test)
# print('Accuracy Score is {:.5}'.format(accuracy_score(y_test, y_predict)))
# print(pd.DataFrame(confusion_matrix(y_test,y_predict)))
```

```
In [15]: # submission = pd.DataFrame()

# submission['ID'] = test_data['ID']
# test_data_temp = test_data[['Gender', 'Age', 'Region_Code', 'Occupation', 'Channel_Code',
#               'Credit_Product', 'Avg_Account_Balance', 'Is_Active', 'begin_month', 'target']]

# submission['Is_Lead'] = model.predict(test_data_temp)
# submission.to_csv('sample_submission.csv')
# submission.head()
```

Showing important features:

```
In [16]: def plot_importance(classifer, x_train, point_size = 25):  
        '''plot feature importance'''  
        values = sorted(zip(x_train.columns, classifer.feature_importances_), key = lambda x: x[1] * -1)  
        imp = pd.DataFrame(values, columns = ["Name", "Score"])  
        imp.sort_values(by = 'Score', inplace = True)  
        sns.scatterplot(x = 'Score', y='Name', linewidth = 0,  
                        data = imp, s = point_size, color='red').set(  
            xlabel='importance',  
            ylabel='features')  
  
        # plot_importance(model, X_train, 20)
```

```
In [17]: # print(model.booster_.feature_importance(importance_type='gain'))
```

## Xgboost

```
In [18]: model = XGBClassifier(max_depth=12,
                               n_estimators=400,
                               min_child_weight=8,
                               subsample=0.8,
                               learning_rate =0.01,
                               seed=42)

eval_set = [(X_test, y_test)]
model.fit(X_train, y_train, early_stopping_rounds=25, eval_metric="auc", eval_set=eval_set, verbose=True)

y_predict = model.predict(X_test)
print('Accuracy Score is {:.5}'.format(accuracy_score(y_test, y_predict)))
print(pd.DataFrame(confusion_matrix(y_test,y_predict)))

plot_importance(model, X_train, 20)
```

```
[0]    validation_0-auc:0.906367
Will train until validation_0-auc hasn't improved in 25 rounds.
[1]    validation_0-auc:0.908923
[2]    validation_0-auc:0.909655
[3]    validation_0-auc:0.909875
[4]    validation_0-auc:0.910589
[5]    validation_0-auc:0.910542
[6]    validation_0-auc:0.910711
[7]    validation_0-auc:0.910832
[8]    validation_0-auc:0.911081
[9]    validation_0-auc:0.911471
[10]   validation_0-auc:0.911584
[11]   validation_0-auc:0.911648
[12]   validation_0-auc:0.911733
[13]   validation_0-auc:0.91202
[14]   validation_0-auc:0.911995
[15]   validation_0-auc:0.911937
[16]   validation_0-auc:0.912135
[17]   validation_0-auc:0.9122
[18]   validation_0-auc:0.912219
[19]   validation_0-auc:0.912248
[20]   validation_0-auc:0.912326
[21]   validation_0-auc:0.912389
[22]   validation_0-auc:0.912448
[23]   validation_0-auc:0.912509
[24]   validation_0-auc:0.912508
[25]   validation_0-auc:0.912541
[26]   validation_0-auc:0.912561
[27]   validation_0-auc:0.912595
[28]   validation_0-auc:0.912642
[29]   validation_0-auc:0.912702
[30]   validation_0-auc:0.912771
[31]   validation_0-auc:0.912855
[32]   validation_0-auc:0.912985
[33]   validation_0-auc:0.913005
[34]   validation_0-auc:0.913002
[35]   validation_0-auc:0.913024
[36]   validation_0-auc:0.913085
[37]   validation_0-auc:0.913128
[38]   validation_0-auc:0.913202
[39]   validation_0-auc:0.913249
```

[40] validation\_0-auc:0.913274  
[41] validation\_0-auc:0.913274  
[42] validation\_0-auc:0.913506  
[43] validation\_0-auc:0.913564  
[44] validation\_0-auc:0.913604  
[45] validation\_0-auc:0.913708  
[46] validation\_0-auc:0.913734  
[47] validation\_0-auc:0.913774  
[48] validation\_0-auc:0.913801  
[49] validation\_0-auc:0.913878  
[50] validation\_0-auc:0.913917  
[51] validation\_0-auc:0.913964  
[52] validation\_0-auc:0.914049  
[53] validation\_0-auc:0.914104  
[54] validation\_0-auc:0.914189  
[55] validation\_0-auc:0.914241  
[56] validation\_0-auc:0.914305  
[57] validation\_0-auc:0.914366  
[58] validation\_0-auc:0.914411  
[59] validation\_0-auc:0.914484  
[60] validation\_0-auc:0.914509  
[61] validation\_0-auc:0.914512  
[62] validation\_0-auc:0.914563  
[63] validation\_0-auc:0.914595  
[64] validation\_0-auc:0.914613  
[65] validation\_0-auc:0.91465  
[66] validation\_0-auc:0.914723  
[67] validation\_0-auc:0.914765  
[68] validation\_0-auc:0.914801  
[69] validation\_0-auc:0.914901  
[70] validation\_0-auc:0.914941  
[71] validation\_0-auc:0.914963  
[72] validation\_0-auc:0.915005  
[73] validation\_0-auc:0.915153  
[74] validation\_0-auc:0.915283  
[75] validation\_0-auc:0.915298  
[76] validation\_0-auc:0.915339  
[77] validation\_0-auc:0.91538  
[78] validation\_0-auc:0.91541  
[79] validation\_0-auc:0.915421  
[80] validation\_0-auc:0.915443  
[81] validation\_0-auc:0.91546

[82] validation\_0-auc:0.915518  
[83] validation\_0-auc:0.915649  
[84] validation\_0-auc:0.915686  
[85] validation\_0-auc:0.91572  
[86] validation\_0-auc:0.915742  
[87] validation\_0-auc:0.91582  
[88] validation\_0-auc:0.915834  
[89] validation\_0-auc:0.915856  
[90] validation\_0-auc:0.915875  
[91] validation\_0-auc:0.915897  
[92] validation\_0-auc:0.916027  
[93] validation\_0-auc:0.916046  
[94] validation\_0-auc:0.916074  
[95] validation\_0-auc:0.916088  
[96] validation\_0-auc:0.916226  
[97] validation\_0-auc:0.916303  
[98] validation\_0-auc:0.916331  
[99] validation\_0-auc:0.916402  
[100] validation\_0-auc:0.916519  
[101] validation\_0-auc:0.916571  
[102] validation\_0-auc:0.916642  
[103] validation\_0-auc:0.916707  
[104] validation\_0-auc:0.916732  
[105] validation\_0-auc:0.916795  
[106] validation\_0-auc:0.916852  
[107] validation\_0-auc:0.916887  
[108] validation\_0-auc:0.91697  
[109] validation\_0-auc:0.917024  
[110] validation\_0-auc:0.917112  
[111] validation\_0-auc:0.9172  
[112] validation\_0-auc:0.917222  
[113] validation\_0-auc:0.917267  
[114] validation\_0-auc:0.917368  
[115] validation\_0-auc:0.917404  
[116] validation\_0-auc:0.917456  
[117] validation\_0-auc:0.91755  
[118] validation\_0-auc:0.917592  
[119] validation\_0-auc:0.917646  
[120] validation\_0-auc:0.917737  
[121] validation\_0-auc:0.917758  
[122] validation\_0-auc:0.917833  
[123] validation\_0-auc:0.917917

[124] validation\_0-auc:0.917992  
[125] validation\_0-auc:0.918056  
[126] validation\_0-auc:0.918097  
[127] validation\_0-auc:0.918128  
[128] validation\_0-auc:0.918198  
[129] validation\_0-auc:0.918244  
[130] validation\_0-auc:0.918302  
[131] validation\_0-auc:0.918334  
[132] validation\_0-auc:0.918357  
[133] validation\_0-auc:0.918453  
[134] validation\_0-auc:0.918483  
[135] validation\_0-auc:0.91854  
[136] validation\_0-auc:0.918584  
[137] validation\_0-auc:0.918625  
[138] validation\_0-auc:0.918664  
[139] validation\_0-auc:0.918701  
[140] validation\_0-auc:0.918711  
[141] validation\_0-auc:0.91876  
[142] validation\_0-auc:0.9188  
[143] validation\_0-auc:0.918847  
[144] validation\_0-auc:0.918864  
[145] validation\_0-auc:0.918887  
[146] validation\_0-auc:0.918934  
[147] validation\_0-auc:0.918969  
[148] validation\_0-auc:0.919  
[149] validation\_0-auc:0.919034  
[150] validation\_0-auc:0.919059  
[151] validation\_0-auc:0.919094  
[152] validation\_0-auc:0.919123  
[153] validation\_0-auc:0.919177  
[154] validation\_0-auc:0.919193  
[155] validation\_0-auc:0.91923  
[156] validation\_0-auc:0.919268  
[157] validation\_0-auc:0.919292  
[158] validation\_0-auc:0.919364  
[159] validation\_0-auc:0.919389  
[160] validation\_0-auc:0.91941  
[161] validation\_0-auc:0.919438  
[162] validation\_0-auc:0.91946  
[163] validation\_0-auc:0.919493  
[164] validation\_0-auc:0.919517  
[165] validation\_0-auc:0.919535

[166] validation\_0-auc:0.919552  
[167] validation\_0-auc:0.919579  
[168] validation\_0-auc:0.919592  
[169] validation\_0-auc:0.919605  
[170] validation\_0-auc:0.919629  
[171] validation\_0-auc:0.919645  
[172] validation\_0-auc:0.919661  
[173] validation\_0-auc:0.919697  
[174] validation\_0-auc:0.919727  
[175] validation\_0-auc:0.919744  
[176] validation\_0-auc:0.91977  
[177] validation\_0-auc:0.919799  
[178] validation\_0-auc:0.919833  
[179] validation\_0-auc:0.919862  
[180] validation\_0-auc:0.919893  
[181] validation\_0-auc:0.919906  
[182] validation\_0-auc:0.919938  
[183] validation\_0-auc:0.919976  
[184] validation\_0-auc:0.919994  
[185] validation\_0-auc:0.920027  
[186] validation\_0-auc:0.920066  
[187] validation\_0-auc:0.920092  
[188] validation\_0-auc:0.920116  
[189] validation\_0-auc:0.920146  
[190] validation\_0-auc:0.920171  
[191] validation\_0-auc:0.920212  
[192] validation\_0-auc:0.920247  
[193] validation\_0-auc:0.920273  
[194] validation\_0-auc:0.92029  
[195] validation\_0-auc:0.92032  
[196] validation\_0-auc:0.920343  
[197] validation\_0-auc:0.920362  
[198] validation\_0-auc:0.920401  
[199] validation\_0-auc:0.92043  
[200] validation\_0-auc:0.920442  
[201] validation\_0-auc:0.920469  
[202] validation\_0-auc:0.92054  
[203] validation\_0-auc:0.920575  
[204] validation\_0-auc:0.920613  
[205] validation\_0-auc:0.920643  
[206] validation\_0-auc:0.920682  
[207] validation\_0-auc:0.920707



[208] validation\_0-auc:0.920737  
[209] validation\_0-auc:0.920767  
[210] validation\_0-auc:0.920806  
[211] validation\_0-auc:0.920832  
[212] validation\_0-auc:0.920875  
[213] validation\_0-auc:0.920886  
[214] validation\_0-auc:0.92092  
[215] validation\_0-auc:0.920969  
[216] validation\_0-auc:0.921015  
[217] validation\_0-auc:0.921065  
[218] validation\_0-auc:0.921125  
[219] validation\_0-auc:0.921147  
[220] validation\_0-auc:0.921178  
[221] validation\_0-auc:0.921205  
[222] validation\_0-auc:0.921243  
[223] validation\_0-auc:0.921273  
[224] validation\_0-auc:0.921326  
[225] validation\_0-auc:0.921341  
[226] validation\_0-auc:0.921358  
[227] validation\_0-auc:0.921417  
[228] validation\_0-auc:0.921437  
[229] validation\_0-auc:0.921469  
[230] validation\_0-auc:0.921522  
[231] validation\_0-auc:0.921552  
[232] validation\_0-auc:0.921599  
[233] validation\_0-auc:0.921608  
[234] validation\_0-auc:0.921654  
[235] validation\_0-auc:0.921669  
[236] validation\_0-auc:0.921705  
[237] validation\_0-auc:0.921713  
[238] validation\_0-auc:0.921726  
[239] validation\_0-auc:0.921747  
[240] validation\_0-auc:0.921759  
[241] validation\_0-auc:0.921813  
[242] validation\_0-auc:0.921834  
[243] validation\_0-auc:0.921853  
[244] validation\_0-auc:0.921899  
[245] validation\_0-auc:0.921932  
[246] validation\_0-auc:0.921979  
[247] validation\_0-auc:0.92202  
[248] validation\_0-auc:0.922056  
[249] validation\_0-auc:0.922076

[250] validation\_0-auc:0.922093  
[251] validation\_0-auc:0.922133  
[252] validation\_0-auc:0.922165  
[253] validation\_0-auc:0.922196  
[254] validation\_0-auc:0.922238  
[255] validation\_0-auc:0.922272  
[256] validation\_0-auc:0.922307  
[257] validation\_0-auc:0.922348  
[258] validation\_0-auc:0.922399  
[259] validation\_0-auc:0.922422  
[260] validation\_0-auc:0.922464  
[261] validation\_0-auc:0.922477  
[262] validation\_0-auc:0.922525  
[263] validation\_0-auc:0.922553  
[264] validation\_0-auc:0.922582  
[265] validation\_0-auc:0.922604  
[266] validation\_0-auc:0.922641  
[267] validation\_0-auc:0.922688  
[268] validation\_0-auc:0.92273  
[269] validation\_0-auc:0.922747  
[270] validation\_0-auc:0.922763  
[271] validation\_0-auc:0.922774  
[272] validation\_0-auc:0.922788  
[273] validation\_0-auc:0.922835  
[274] validation\_0-auc:0.922855  
[275] validation\_0-auc:0.922876  
[276] validation\_0-auc:0.922891  
[277] validation\_0-auc:0.922909  
[278] validation\_0-auc:0.922949  
[279] validation\_0-auc:0.922972  
[280] validation\_0-auc:0.92299  
[281] validation\_0-auc:0.923029  
[282] validation\_0-auc:0.923042  
[283] validation\_0-auc:0.923058  
[284] validation\_0-auc:0.923102  
[285] validation\_0-auc:0.923138  
[286] validation\_0-auc:0.923181  
[287] validation\_0-auc:0.9232  
[288] validation\_0-auc:0.923215  
[289] validation\_0-auc:0.923245  
[290] validation\_0-auc:0.923287  
[291] validation\_0-auc:0.92332

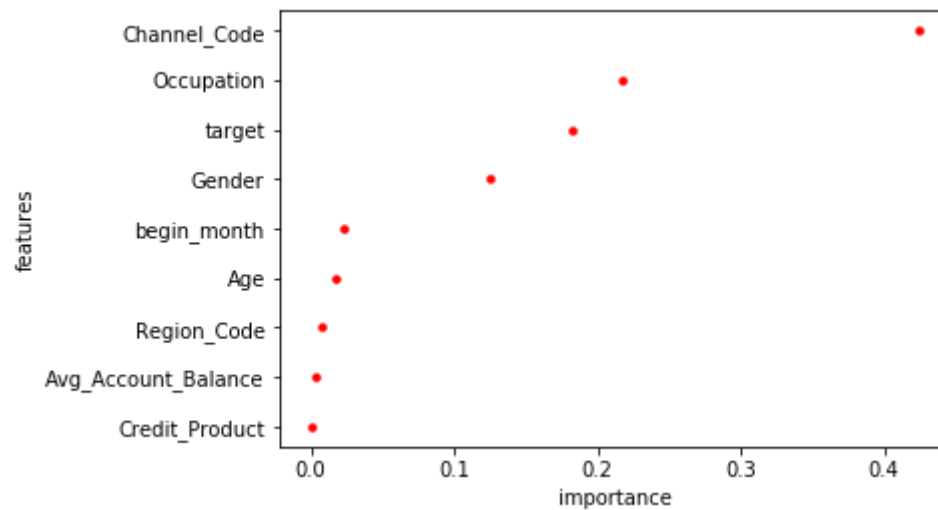
[292] validation\_0-auc:0.92335  
[293] validation\_0-auc:0.923386  
[294] validation\_0-auc:0.923433  
[295] validation\_0-auc:0.923463  
[296] validation\_0-auc:0.923498  
[297] validation\_0-auc:0.923529  
[298] validation\_0-auc:0.923553  
[299] validation\_0-auc:0.923563  
[300] validation\_0-auc:0.923592  
[301] validation\_0-auc:0.923614  
[302] validation\_0-auc:0.923631  
[303] validation\_0-auc:0.923649  
[304] validation\_0-auc:0.923686  
[305] validation\_0-auc:0.923706  
[306] validation\_0-auc:0.923725  
[307] validation\_0-auc:0.923753  
[308] validation\_0-auc:0.923778  
[309] validation\_0-auc:0.923796  
[310] validation\_0-auc:0.923824  
[311] validation\_0-auc:0.923848  
[312] validation\_0-auc:0.923861  
[313] validation\_0-auc:0.92389  
[314] validation\_0-auc:0.923925  
[315] validation\_0-auc:0.923946  
[316] validation\_0-auc:0.923964  
[317] validation\_0-auc:0.923988  
[318] validation\_0-auc:0.924016  
[319] validation\_0-auc:0.924026  
[320] validation\_0-auc:0.924039  
[321] validation\_0-auc:0.924071  
[322] validation\_0-auc:0.924088  
[323] validation\_0-auc:0.924117  
[324] validation\_0-auc:0.924151  
[325] validation\_0-auc:0.924177  
[326] validation\_0-auc:0.924205  
[327] validation\_0-auc:0.924214  
[328] validation\_0-auc:0.924227  
[329] validation\_0-auc:0.924257  
[330] validation\_0-auc:0.924287  
[331] validation\_0-auc:0.924306  
[332] validation\_0-auc:0.924336  
[333] validation\_0-auc:0.924367

[334] validation\_0-auc:0.924394  
[335] validation\_0-auc:0.924419  
[336] validation\_0-auc:0.924442  
[337] validation\_0-auc:0.92446  
[338] validation\_0-auc:0.924494  
[339] validation\_0-auc:0.924508  
[340] validation\_0-auc:0.924524  
[341] validation\_0-auc:0.924545  
[342] validation\_0-auc:0.924575  
[343] validation\_0-auc:0.924601  
[344] validation\_0-auc:0.924622  
[345] validation\_0-auc:0.924648  
[346] validation\_0-auc:0.924671  
[347] validation\_0-auc:0.92469  
[348] validation\_0-auc:0.92471  
[349] validation\_0-auc:0.924731  
[350] validation\_0-auc:0.924763  
[351] validation\_0-auc:0.924786  
[352] validation\_0-auc:0.924815  
[353] validation\_0-auc:0.924836  
[354] validation\_0-auc:0.92485  
[355] validation\_0-auc:0.924871  
[356] validation\_0-auc:0.9249  
[357] validation\_0-auc:0.924917  
[358] validation\_0-auc:0.924933  
[359] validation\_0-auc:0.924946  
[360] validation\_0-auc:0.924964  
[361] validation\_0-auc:0.924982  
[362] validation\_0-auc:0.924996  
[363] validation\_0-auc:0.925013  
[364] validation\_0-auc:0.925034  
[365] validation\_0-auc:0.925068  
[366] validation\_0-auc:0.925091  
[367] validation\_0-auc:0.925124  
[368] validation\_0-auc:0.925156  
[369] validation\_0-auc:0.925176  
[370] validation\_0-auc:0.925186  
[371] validation\_0-auc:0.925212  
[372] validation\_0-auc:0.925231  
[373] validation\_0-auc:0.925258  
[374] validation\_0-auc:0.925276  
[375] validation\_0-auc:0.925298

```
[376] validation_0-auc:0.925324
[377] validation_0-auc:0.925361
[378] validation_0-auc:0.925389
[379] validation_0-auc:0.925392
[380] validation_0-auc:0.925406
[381] validation_0-auc:0.925437
[382] validation_0-auc:0.925461
[383] validation_0-auc:0.925484
[384] validation_0-auc:0.925506
[385] validation_0-auc:0.925523
[386] validation_0-auc:0.925544
[387] validation_0-auc:0.925571
[388] validation_0-auc:0.925585
[389] validation_0-auc:0.925598
[390] validation_0-auc:0.925613
[391] validation_0-auc:0.925626
[392] validation_0-auc:0.925649
[393] validation_0-auc:0.92567
[394] validation_0-auc:0.925688
[395] validation_0-auc:0.925698
[396] validation_0-auc:0.925727
[397] validation_0-auc:0.925736
[398] validation_0-auc:0.925755
[399] validation_0-auc:0.925772
```

Accuracy Score is 0.84758

	0	1
0	34971	2429
1	8999	28576



## CatBoost

```
In [19]: # model = CatBoostClassifier(iterations=250,
#                                   learning_rate=0.2,
#                                   od_type='Iter',
#                                   verbose=25,
#                                   depth=16,
#                                   random_seed=42)

# model.fit(X_train, y_train)
# y_predict = model.predict(X_test)
# print('CatBoost Accuracy Score is {:.5}'.format(accuracy_score(y_test, y_predict)))
# print(pd.DataFrame(confusion_matrix(y_test, y_predict)))
```

```
>>> roc_auc_score(y, clf.predict_proba(X)[:, 1])
```

```
In [24]: submission = pd.DataFrame()

submission['ID'] = test_data['ID']
test_data_temp = test_data[['Gender', 'Age', 'Region_Code', 'Occupation', 'Channel_Code',
                             'Credit_Product', 'Avg_Account_Balance', 'target', 'begin_month']]

submission['Is_Lead'] = model.predict(test_data_temp)

submission.to_csv('sample_submission.csv', index=False)
submission.head()
```

Out[24]:

	ID	Is_Lead
0	VBENBARO	0
1	CCMEWNKY	0
2	VK3KGA9M	0
3	TT8RPZVC	0
4	SHQZEY TZ	0

## Building the Keras neural networks

After a good deal of trial and error, I found that a network architecture with three hidden layers, each followed by a dropout layer of rate 0.3, was as good as I could find. I used ReLU activation in those hidden layers, and adam optimization and a loss metric of mean squared error in the model as a whole. I also settled on a mean squared logarithmic error loss function, since it performed better than mean absolute error, mean squared error, and mean absolute percentage error.

The dataset being so large, I had great results increasing the batch size for the first couple models.

```
In [21]: # from sklearn.model_selection import train_test_split
# from sklearn_pandas import DataFrameMapper
# from sklearn.preprocessing import OneHotEncoder, OrdinalEncoder, StandardScaler
# from tensorflow.keras import Sequential, Input
# from tensorflow.keras.layers import Dense, Dropout
# from tensorflow.keras.callbacks import EarlyStopping

# def run_pipeline(batch_size):

#     input_nodes = X_train.shape[1]
#     output_nodes = 1

#     model = Sequential()
#     model.add(Input((input_nodes,)))
#     model.add(Dense(32, activation="sigmoid"))
#     model.add(Dropout(0.4, seed=0))
#     model.add(Dense(16, activation="sigmoid"))
#     model.add(Dropout(0.4, seed=1))
#     model.add(Dense(8, activation="sigmoid"))
#     model.add(Dropout(0.4, seed=2))
#     model.add(Dense(output_nodes, activation='sigmoid'))
#     model.compile(optimizer="adam", loss="mean_squared_error", metrics=['accuracy'])

#     es = EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=40)

#     history = model.fit(
#         X_train,
#         y_train,
#         batch_size=batch_size,
#         epochs=500,
#         validation_data=(X_test, y_test),
#         verbose=2,
#         callbacks=[es]
#     )

#     return history.history, model
```



```
In [22]: # print("Model 1:")  
# history_1, model1 = run_pipeline(batch_size=1024*10)
```

```
In [23]: # submission = pd.DataFrame()  
  
# submission['ID'] = test_data['ID']  
# test_data_temp = test_data[['Gender', 'Age', 'Region_Code', 'Occupation', 'Channel_Code',  
#                             'Credit_Product', 'Avg_Account_Balance', 'target', 'begin_month']]  
  
# submission['Is_Lead'] = model1.predict(test_data_temp)  
  
# submission.to_csv('sample_submission.csv', index=False)  
# submission.head()
```